

Introduction to Deep Learning

Assignment 2 Report

Tom Mahler, student ID: 204639926
Chang Liu, student ID: 850259201

June 9, 2019

1 Introduction

In this assignment, we use convolutional neural networks (CNNs) to carry out the task of facial recognition, since CNNs are the current state-of-the-art approach for analyzing image based datasets. More specifically, we implement a one-shot classification solution. Wikipedia defines one-shot learning as follows:

“... an object categorization problem, found mostly in computer vision. Whereas most machine learning based object categorization algorithms require training on hundreds or thousands of samples/images and very large datasets, one-shot learning aims to learn information about object categories from one, or only a few, training samples/images.”

Our work is based on the paper [Siamese Neural Networks for One-shot Image Recognition](#). Our goal, like that of the paper, is to successfully execute a one-shot learning task for previously unseen objects. Given two facial images of previously unseen persons, our architecture will have to successfully determine whether they are the same person. While we are encouraged to use the architecture described in this paper as a starting point, we shall explore other possibilities as well.

In this work, we are making use of [TensorFlow 2.0](#) API on Google Colab notebook set on GPU runtime (GPU: 1xTesla K80, having 2496 CUDA cores, compute 3.7, 12GB (11.439GB Usable) GDDR5 VRAM).

We load TF2.0 as follows:

```
[0]: %%capture
from __future__ import absolute_import, division, print_function,
    ↳ unicode_literals

!pip install -q tensorflow-gpu==2.0.0-alpha0
import tensorflow as tf
# AUTOTUNE = tf.data.experimental.AUTOTUNE
```

For reproducible results, we init numpy and tensorflow random seed with same values

```
[0]: import numpy as np
np.set_printoptions(precision=2)
```

```

from numpy.random import seed
from tensorflow.random import set_seed
seed(2)
set_seed(2)

```

```

[3]: print('Running on GPU' if tf.test.is_gpu_available() else 'Please change runtime_
      ↳type to GPU on Google Colab under Runtime') # Make sure we are set to GPU_
      ↳(under Runtime->Change runtime type)

```

Running on GPU

2 Organizing the Data

We use the [Labeled Faces in the Wild](#) dataset. Note that there are several versions of this dataset, we use the version found [here](#) (it's called LFW-a, and is also used in the DeepFace paper).

We use the following train and test sets to train your model: [Train Test](#). We use the test set to perform one-shot learning. This division is set up so that no subject from the test set is included in the train set.

The downloaded directory is built in the following structure:

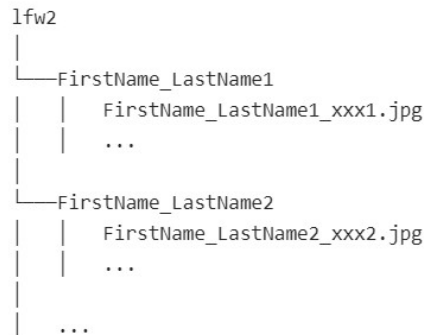


Figure 1: Downloaded directory structure

Which includes both the train, test, and additional data. Therefore, our first task was to organize the data better. We chose to use python scripts to create 3 flattened directories: test, train, and unused. For this task, we cannot use additional data for the training, however, we maintain the unused directory for potential future work.

Note that you first must download the dataset manually

2.1 Organize dataset

This section explains how we organized the dataset. **Note that the organized dataset was uploaded to our [GitHub account: MahlerTom/Siamese-Neural-Networks](#), and thus this section can be skipped**

First we download all the necessary files: `pairsDevTrain.txt`, `pairsDevTest.txt`, and `lfw2.zip`, and we unzip the dataset

```
[0]: %%capture
!wget http://vis-www.cs.umass.edu/lfw/pairsDevTrain.txt
!wget http://vis-www.cs.umass.edu/lfw/pairsDevTest.txt
!gdown https://drive.google.com/uc?id=1p1wjaqpTh_5RHfJu4vUh8JJCdKwYMHCP
!unzip lfw.zip
```

Next, we define several functions that help us:

1. The `create_pairs` function helps us create a list of all the names that appear in the `pairsDevTrain.txt` and `pairsDevTest.txt` which we will later use for moving the folders.
2. The `flatten` function simply flatten a given folder with the specific structure mentioned.
3. The `move_dirs_and_flatten` function moves a folder from the `src_path` using the `folder_names` list that was created using `create_pairs` into a new folder inside `dst_path` and then flatten it using `flatten`.

```
[0]: import shutil
import os

def create_pairs(pairs_path):
    names = set()
    with open(pairs_path) as pairs_path_f:
        pairs_list = pairs_path_f.readlines()[1:]

    for pair in pairs_list:
        pair = pair[:-1].split('\t')
        if len(pair) == 3:
            names.add(pair[0])
        elif len(pair) == 4:
            names.add(pair[0])
            names.add(pair[2])
    return list(names)

def flatten(src, verbose=0):
    for directory in os.listdir(src):
        for file in os.listdir(src + directory):
            if verbose > 0:
                print("Moving " + file + "...")
            shutil.move(src + directory + '/' + file, src + file)

def move_dirs_and_flatten(src_path, dst_path, folders_names, verbose=0):
    for folder_name in folders_names:
        if verbose > 0:
            print("Moving " + folder_name + "...")
        shutil.move(src_path + folder_name, dst_path + folder_name)

    flatten(dst_path)
```

With these functions, we organize our dataset into a new folder with the following structure:

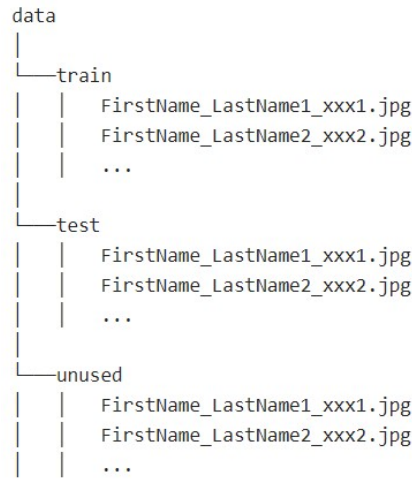


Figure 2: Organized directory structure

```
[0]: import shutil

train_names = create_pairs('pairsDevTrain.txt')
test_names = create_pairs('pairsDevTest.txt')

src = 'lfw2/lfw2/'
move_dirs_and_flatten(src, 'data/train/', train_names)
move_dirs_and_flatten(src, 'data/test/', test_names)
flatten(src)
shutil.move(src, 'data/unused/')
!rm -r lfw2
```

3 Installation

The dataset was uploaded to [MahlerTom/Siamese-Neural-Networks](#), so we first need to clone the repository, with the data. To make things easier, we also define:

- `repo_path` – the repository path (this should be cross platform since we use `os` module).
- `train_path` – the train dataset path.
- `test_path` – the test dataset path.

```
[0]: %%capture
import os

# Clone the entire repo.
!git clone -s git://github.com/MahlerTom/Siamese-Neural-Networks.git
↪SiameseNeuralNetworks
repo_path = os.path.join(os.getcwd(), 'SiameseNeuralNetworks')
```

```
train_path = os.path.join(repo_path, 'data', 'train')
test_path = os.path.join(repo_path, 'data', 'test')
```

4 Analysis of the Dataset

Before we begin our training, we need to prepare the dataset. Since we are using TensorFlow 2.0, we will make use of its functions. We followed the guide at: https://www.tensorflow.org/alpha/tutorials/load_data/images

4.1 Preparing the dataset

Before we begin our training, we need to prepare the dataset. Since we are using TensorFlow 2.0, we will make use of its functions. We followed the guide at: https://www.tensorflow.org/alpha/tutorials/load_data/images

Tensorflow makes use of smart functions that can load images given their paths. In addition, we received trainPairs.txt and testPairs.txt, which include the labels.

The data structure is as follows:

```
((left_img_path, right_img_path), label)
```

Thus, the following load_data function will create it.

```
[8]: from SiameseNeuralNetworks.data import load_data

# Utility array to print the name of the label easily (tested on Google Colab)
label_name = ['Different', 'Same']

train_paths_labels = load_data(train_path, labels_file=os.path.join(repo_path,
    → 'trainPairs.txt'), print_imgs=3)
test_paths_labels = load_data(test_path, labels_file=os.path.join(repo_path,
    → 'testPairs.txt'), print_imgs=3)
```

Loaded 6685 image paths

#####

Printing Example Images



(a) Lleyton Hewitt 0033



(b) Kim Clijsters 0004



(c) Roh Moo-hyun 0020

```
#####
Loaded 2741 image paths
#####
Printing Example Images
```



(a) Richard Gere 0007



(b) Nancy Pelosi 0002



(c) Allison Searing 0001

The images are grayscale with resolution 250×250 .

```
[9]: from SiameseNeuralNetworks.utils import print_dataset_stat

print_dataset_stat(train_paths_labels, ds_name='Train')
print('#####\n\n#####')
print_dataset_stat(test_paths_labels, ds_name='Test')
```

Examples for Train dataset structure:

#####

Raw view:

```
[('/content/SiameseNeuralNetworks/data/train/Aaron_Peirsol_0001.jpg',
'/content/SiameseNeuralNetworks/data/train/Aaron_Peirsol_0002.jpg', 1),
('/content/SiameseNeuralNetworks/data/train/Aaron_Peirsol_0003.jpg',
'/content/SiameseNeuralNetworks/data/train/Aaron_Peirsol_0004.jpg', 1),
('/content/SiameseNeuralNetworks/data/train/Aaron_Sorkin_0001.jpg',
'/content/SiameseNeuralNetworks/data/train/Aaron_Sorkin_0002.jpg', 1)]
```

#####

Pretty view:

```
[('Aaron_Peirsol_0001', 'Aaron_Peirsol_0002', 'Same'), ('Aaron_Peirsol_0003',
'Aaron_Peirsol_0004', 'Same'), ('Aaron_Sorkin_0001', 'Aaron_Sorkin_0002',
'Same')]
```

#####

Dataset size: 2200

2 classes:

Class	Size	Percentage
Different	1100/2200	50.00%
Same	1100/2200	50.00%

```

+-----+-----+-----+
#####

#####
Examples for Test dataset structure:
#####
Raw view:
(['/content/SiameseNeuralNetworks/data/test/Abdullah_Gul_0013.jpg',
'/content/SiameseNeuralNetworks/data/test/Abdullah_Gul_0014.jpg', 1),
('/content/SiameseNeuralNetworks/data/test/Abdullah_Gul_0013.jpg',
'/content/SiameseNeuralNetworks/data/test/Abdullah_Gul_0016.jpg', 1),
('/content/SiameseNeuralNetworks/data/test/Abdullatif_Sener_0001.jpg',
'/content/SiameseNeuralNetworks/data/test/Abdullatif_Sener_0002.jpg', 1)]
#####
Pretty view:
(['Abdullah_Gul_0013', 'Abdullah_Gul_0014', 'Same'), ('Abdullah_Gul_0013',
'Abdullah_Gul_0016', 'Same'), ('Abdullatif_Sener_0001', 'Abdullatif_Sener_0002',
'Same')]
#####
Dataset size: 1000
2 classes:
+-----+-----+-----+
| Class | Size | Percentage |
+-----+-----+-----+
| Different | 500/1000 | 50.00% |
| Same | 500/1000 | 50.00% |
+-----+-----+-----+

```

4.2 Create Validation Set

It is often recommended to use a validation set when training the model; thus, we have chosen to use a 70-30 train-validation ratio. It's important to note that while we are guaranteed that there is no subject from the test set included in the train set, we do not have this guarantee for the train-validation separation. Therefore, if we randomly choose a 70-30 separation ratio (a common choice), we will likely have subjects included in both the training and validation sets, which will cause problems.

Therefore, we explore separating the sets based on the person's name, to guarantee the sets are independent.

As we can see, if we choose the letters: A, B, C, D, E for the *same* validation set, and the letters: A, B, C for the *different* validation set, we will get:

```

[10]: from SiameseNeuralNetworks.utils import explore_subject_names
      from SiameseNeuralNetworks.data import image_name

      from prettytable import PrettyTable

      train_same = [image_name(p[0]) for p in train_paths_labels[:1100]]
      train_diff = [image_name(p[0]) for p in train_paths_labels[1100:]]

```

```

h_same = explore_subject_names(train_same)
h_diff = explore_subject_names(train_diff)

t = PrettyTable([''] + list(h_same.keys()))
t.add_row(['Same'] + list(h_same.values()))
t.add_row(['Different'] + list(h_diff.values()))
print(t)
print()

letters_selected_same = h_same['A'] + h_same['B'] + h_same['C'] + h_same['D'] +
    h_same['E']
letters_selected_different = h_diff['A'] + h_diff['B'] + h_diff['C']

t = PrettyTable(['', 'Total Size', 'Chosen Letters', 'Chosen Letters Size',
    'Percentage'])
t.add_row(['Same', sum(h_same.values()), 'A, B, C, D, E', letters_selected_same,
    f'{letters_selected_same/len(train_same)*100:.2f}%'])
t.add_row(['Different', sum(h_diff.values()), 'A, B, C',
    letters_selected_different, f'{letters_selected_different/len(train_diff)*100:.2f}%'])
t.add_row(['All', sum(h_same.values()) + sum(h_diff.values()), 'A-Z',
    letters_selected_same + letters_selected_different, f'{(letters_selected_same +
    letters_selected_different)/(len(train_same) + len(train_diff))*100:.2f}%'])
print(t)

```

Table 1: Default values used in the experiments

	A	B	C	D	E	F	G	H	I	J	K	L		
Same	101	59	80	58	45	20	50	42	9	145	29	41		
Different	144	121	114	103	64	31	62	47	16	151	34	40		
	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
102	25	4	46	7	74	62	45	1	16	17	3	13	6	
63	22	1	34	2	20	21	8	0	1	1	0	0	0	

Table 2: Chosen train-validation split

	Total Size	Chosen Letters	Chosen Letters Size	Percentage
Same	1100	A, B, C, D, E	343	31.18%
Different	1100	A, B, C	379	34.45%
All	2200	A-Z	722	32.82%

Which provides a good separation. The following function splits the train to train and validation based on the chosen letters:


```
[11]: from SiameseNeuralNetworks.data import split_train_val_paths

train_same, val_same = split_train_val_paths(train_paths_labels[:1100],
↳letters='ABCDE')
train_diff, val_diff = split_train_val_paths(train_paths_labels[1100:],
↳letters='ABC')
train_paths_labels_split = train_same + train_diff
val_paths_labels_split = val_same + val_diff

t = PrettyTable(['', 'Size', 'Percentage'])
t.add_row(['Train Same', f'{len(train_same)}/{len(train_paths_labels)}',
↳f'{len(train_same)/len(train_paths_labels)*100:.2f}%'])
t.add_row(['Train Different', f'{len(train_diff)}/{len(train_paths_labels)}',
↳f'{len(train_diff)/len(train_paths_labels)*100:.2f}%'])
t.add_row(['Train', f'{len(train_paths_labels_split)}/
↳{len(train_paths_labels)}', f'{len(train_paths_labels_split)/
↳len(train_paths_labels)*100:.2f}%'])
t.add_row(['', '', ''])
t.add_row(['Validation Same', f'{len(val_same)}/{len(train_paths_labels)}',
↳f'{len(val_same)/len(train_paths_labels)*100:.2f}%'])
t.add_row(['Validation Different', f'{len(val_diff)}/{len(train_paths_labels)}',
↳f'{len(val_diff)/len(train_paths_labels)*100:.2f}%'])
t.add_row(['Validation', f'{len(val_paths_labels_split)}/
↳{len(train_paths_labels)}', f'{len(val_paths_labels_split)/
↳len(train_paths_labels)*100:.2f}%'])
print(t)
```

	Size	Percentage
Train Same	757/2200	34.41%
Train Different	721/2200	32.77%
Train	1478/2200	67.18%
Validation Same	343/2200	15.59%
Validation Different	379/2200	17.23%
Validation	722/2200	32.82%

- A training set of 757 *same* and 721 *different*, and a total of 1478/2200 = 67.1%.
- A validation set of 343 *same* and 379 *different*, and a total of 722/2200 = 32.8%.

The reason for this specific choice, is that while we are guaranteed that there is no subject from the test set included in the train set, we do not have this guarantee for the train-validation separation. Therefore, if we randomly choose a 70-30 train-validation separation (a common choice), we will likely have validation subjects included in the training set, which will cause a problems.

5 Methods

5.1 Network Architecture

We implemented the Siamese neural network for this task. Before implementing the network, we resized the input images from size of 250×250 to 125×125 .

We start with the exactly same network architecture as in [this paper](#). Generally, the network is composed of two identical convolutional neural networks, which are connected to one fully connected layers. Each CNN is composed of 4 convolutional layers, each of which uses a single channel filters with various size and fixed stride of 1. The network applies a ReLU activation function to the output feature maps, followed by maxpooling with a filter size and stride of 2. The units in the final convolutional layer are flattened into a single vector. This convolutional layer is followed by a fully-connected layer, and then one more layer computing the L1 distance metric between the output of both Siamese twins, which is given to a single sigmoidal output unit.

The original network architecture performs poorly on our task. Assuming that the reason was too many parameters, we reduced the number of filters and the adjusted filter sizes. We tried different network settings and made our decision based on the performance on validation dataset, attempts including:

1. Changing number of filters and filter size:
 - (a) Number of filters: 4, 8, 16
 - (b) Filter size: 2×2 to 12×12 .
2. Changing size of the fully connected layers: 1024 to 4096
3. Weight initialization method
 - (a) $N(0, 10^{-2})$
 - (b) Glorot(Xavier) normal
 - (c) LeCun normal
4. Loss function
 - (a) Binary cross-entropy
 - (b) Contrastive loss
5. Dropout: we try to apply dropout between the layers in a decreasing way. After the first layer we add dropout with rate X and then $X - 0.05$, etc. and in the last layer we apply dropout rate of 0.1.

Changing the loss function to contrastive loss requires some modifications in the network architecture. According to the loss function:

$$Y \frac{1}{2} (D_w)^2 + (1 - Y) \frac{1}{2} \{ \max(\text{margin} - D_w, 0) \}^2 \quad (1)$$

where $Y = 1$ for true class (images are of same people) and 0 vice versa, and D_w is expected to be the euclidean distance of the two feature maps of the input image pairs, we remove the final sigmoidal layer of the network, changing the distance measurement from L1 to euclidean, and used the distance directly as the output of our network.

Eventually, we achieved our best results with the following network architecture:

1. Number of CNN layers: 4
2. Number of filters: 8, 16, 32, 64
3. Filter size: 10×10 , 7×7 , 4×4 , 4×4
4. Size of fully connected layer: 2048
5. Weight initialization method: normal distribution $N(0, 10^{-2})$
6. Bias initialization method: normal distribution $N(0.5, 10^{-2})$
7. Loss function: binary cross-entropy

5.2 Evaluation Methods

In this section we explain the evaluation methods we shall use to measure the performance of our architecture.

5.2.1 Hyper-parameters tuning

We use hyper-parameters tuning in order to choose the best configuration for our network. There are numerous parameters that affect the performance of the network, and we shall focus on the following parameters: resizing the images, batch size, number of filters, number of units in the last layer, optimizer, and learning rate.

As we explained in the methods section. In each layer convolutional layer, we define a different number of filters, so that the next layer has twice the number of filters as the previous layer. The original network includes 64, 128, 256, and 512 filters in the convolutional layers, and 4096 units in the last layer. This results in a very large and deep network, that takes too long to run (and doesn't necessarily output better results). Therefore, we define a configurable filter size, while keeping the proportions between the layers, and a configurable unit size for the last layer. As part of our evaluation, we shall try different filter and unit size to check the effect on the performance.

In addition, we measure the effects of resizing the image. Intuitively, higher resolution (250×250) should result in more details, making the network learn deeper insights. However, for many tasks smaller resolutions are good enough. We will try to resize the network to 125×125 and see test the performance.

Finally, the learning rate, batch size, and optimizer, has a significant effect on the training, so we shall test them as well.

5.2.2 Early Stopping and Model Checkpoint

To avoid over-fitting, and running for too many epochs, we use early stopping. Early stopping means that we train the model while monitoring the validation loss. When we see that the validation loss stops decreasing we stop the training process. We apply patience of three, which means that we monitor the loss and stop only when it hasn't decreased for at least three epochs. The model checkpoint callback helps us save only the best weights even of the loss increased.

5.2.3 Loss and other metrics

We use binary cross-entropy loss and the contrastive loss while monitoring the accuracy. We note that, as we explained in the data section, we split the training set into 70% training set and 30% validation set. In all our experiments, we evaluate our model on the validation and the test set.

6 Experiments and Results

The following tables demonstrate our detailed experiment settings and the corresponding results. We conducted experiments to examine the effect of learning rate, batch size and optimizers (Table 3), weight initialization method (Table 4) and drop-out (Table 6), and corresponding results are presented. Table 5 presents the result when we switch the loss function from binary cross-entropy to contrastive loss. In Figure 5 and Figure 6 we present example graphs of on of the highest accuracy and loss results on train/validation (red/green) dataset during the experiment. We note that we got similar shapes with different parameter choices.




Figure 5: Example of train/val accuracy




Figure 6: Example of train/val loss

By trying different hyper-parameters, we observed some relationships between the hyper-parameters and the performance (under the current network architecture), as shown in fig. 7. This graph provides a straightforward interpretation of the results tables. It can be observed that less filters, smaller learning rate, modest batch size, higher dimension for the final fully-connect layer are preferred choices according to the accuracy. Meanwhile, Adam optimizer works better than SGD and RMS in most cases.




Figure 7: Performance with different hyper-parameters

Table 3: Detailed experiment settings and results with normal distribution weight initialization

Exp No.	Batch	LR	Optimizer	Time	T. Loss	T. Acc	V. Loss	V. Acc
Exp 1.1	32	0.0001	Adam	274.99	46.443	64.80%	46.485	65.37%
Exp 1.2			RMS	273.85	44.981	64.20%	44.934	64.54%
Exp 1.3			SGD	275.91	46.375	64.60%	46.545	62.47%
Exp 2.1	128	0.0001	Adam	250.59	19.7251	66.10%	19.7176	67.04%
Exp 2.2			RMS	251.98	19.7577	67.50%	19.7525	64.54%
Exp 2.3			SGD	252.34	19.7638	65.10%	19.7447	64.82%
Exp 3.1	256	0.0001	Adam	236.72	29.3041	67.50%	29.3122	65.37%
Exp 3.2			RMS	235.97	29.3125	66.20%	29.3159	65.79%
Exp 3.3			SGD	238.04	29.3167	68.30%	29.3146	64.96%
Exp 4.1	32	0.0005	Adam	137.86	9.109	65.00%	0.9306	61.91%
Exp 4.2			RMS	152.09	8.499	67.70%	0.9365	62.47%
Exp 4.3			SGD	129.36	9.702	64.20%	1.0166	61.36%
Exp 5.1	128	0.001	Adam	179.71	1.082	68.40%	1.155	64.54%
Exp 5.2			RMS	167.02	1.119	64.40%	1.173	63.99%
Exp 5.3			SGD	166.05	1.128	64.10%	1.187	62.19%

Table 4: Detailed experiment settings and results with Glorot (Xavier) weight initialization

Exp No.	Batch	LR	Optimizer	Time	T. Loss	T. Acc	V. Loss	V. Acc
Exp 6.1	128	0.0001	Adam	506.69	10.890	63.70%	10.86	64.96%
Exp 6.2		0.0005		333.91	1.980	63.80%	1.969	64.27%
Exp 6.3		0.005		140.10	0.879	50.00%	0.880	47.51%
Exp 6.4		0.01		44.64	1.195	50.00%	1.195	47.51%

Table 5: Detailed experiment settings and results with contrastive loss

Exp No.	Batch	LR	Optimizer	Time	T. Loss	T. Acc	V. Loss	V. Acc
Exp 7.1	128	0.001	Adam	248.47	0.338	65.8%	0.355	61.08%

Table 6: Detailed experiment settings with Drop out

Exp No.	Drop Rate	Batch	LR	Optimizer	Time	T. Loss	T. Acc	V. Loss	V. Acc
Exp 8.1	0.2	128	0.001	Adam	248.47	0.338	65.8%	0.355	61.08%
Exp 8.2	0.25				149.89	1.156	68.60%	1.240	62.88%

Table 7: Detailed experiment settings and results with changing filters and units

Exp No.	Units	Filters	Batch	LR	Optimizer	Time	T. Loss	T. Acc	V. Loss	V. Acc
Exp 9.1	1024	4	128	0.001	Adam	231.57	1.000	65.50%	1.000	64.27%
Exp 9.2		16				357.06	1.023	66.70%	1.073	63.16%
Exp 9.3		32				585.85	0.886	50.00%	0.887	47.51%
Exp 9.4		64				1155.79	1.159	50.00%	1.161	47.51%
Exp 10.1	2048	4	128	0.001	Adam	224.64	0.931	67.70%	0.946	63.57%
Exp 10.2		16				292.79	1.139	67.60%	1.209	63.71%
Exp 10.3		32				585.04	0.792	50.00%	0.794	47.51%
Exp 10.4		64				1153.48	1.065	50.00%	1.067	47.51%

7 Conclusion

In all of our results, we could not reach accuracy higher than 70% (see Figure 8), although we manage to minimize the loss (see Figure 9). It is obvious that the obtained results are not satisfied, but we cannot get any significant improvement with any of the presented configurations, even when changing the loss function (see Table 5) or when changing the network (see Table 6 and Table 7). A possible reason is that the network is weak, meaning that it is not capable of generalizing to unseen image pairs. That's why we achieved almost 100% accuracy on training data but always fail to exceed 70% on validation set, while both losses are decreasing simultaneously. Changing the parameters does not contribute significantly to the performance, and we think it is necessary to reorganize the data, i.e., to change the ratio between same and different class within each mini-batch. Unfortunately we didn't have enough time for that this time, but it will be an interesting attempt to make in the future.

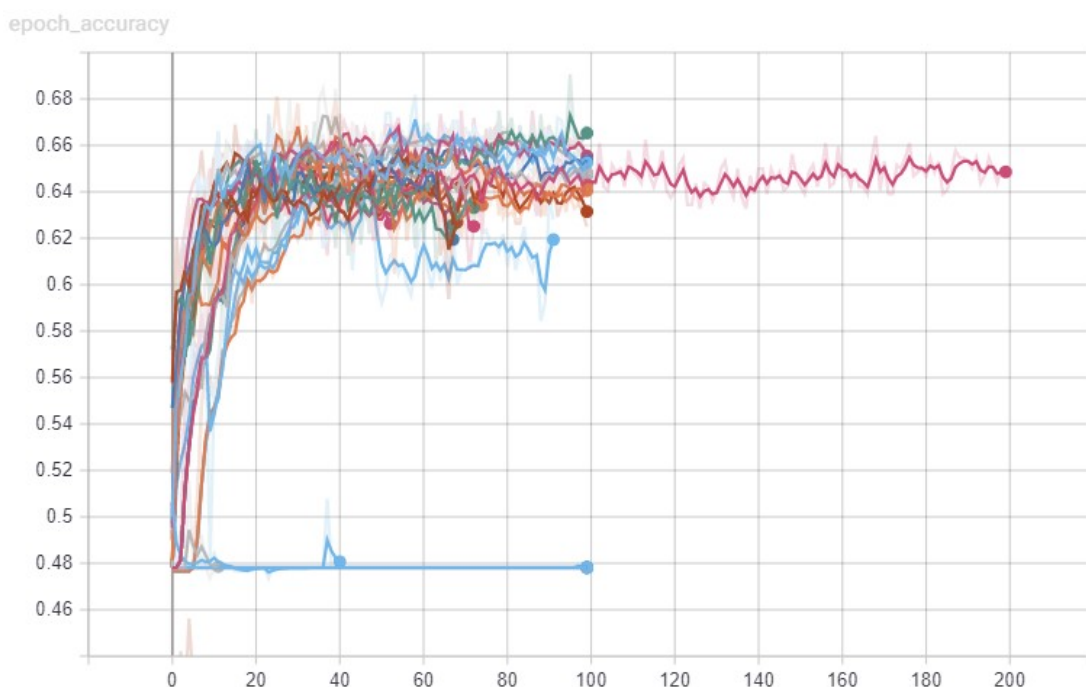


Figure 8: Validation accuracy in each epoch in each experiment

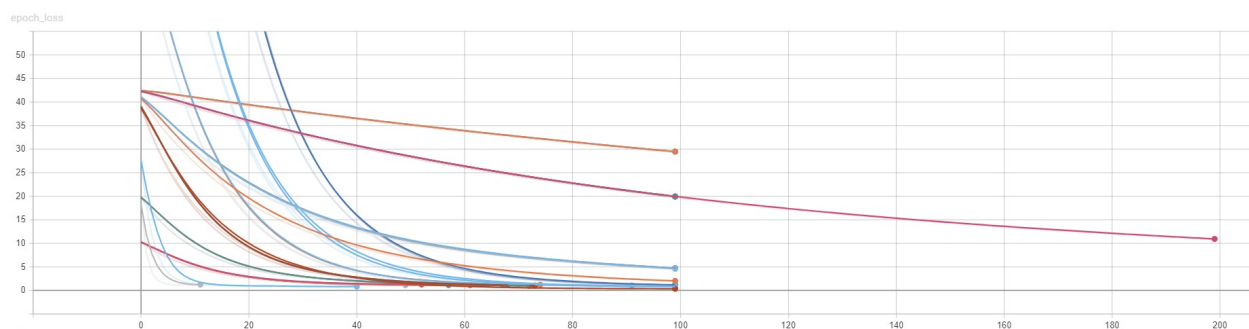


Figure 9: Validation loss in each epoch in each experiment